

# Intro to Machine Learning

EE379K - Architectures for Big Data Sciences

Dr. Sriram Vishwanath

Department of Electrical and Computer Engineering

The University of Texas at Austin

Spring 2016

# What is Machine Learning?

- Extremely broad field that is expanding very fast
- Many different definitions, but in general:
- “[Machine Learning is the] field of study that gives computers the ability to learn without being explicitly programmed.”
  - Arthur Samuel (1959)

# What is Machine Learning?

- Another definition:
- “A computer program is said to learn from experience **E** with respect to some task **T** and some performance measure **P**, if its performance on **T**, as measured by **P**, improves with experience **E**”
  - Tom Mitchell

# Example Tasks

- Predict traffic patterns at an intersection
- Predict stock prices
- Determine if a tumor is cancerous or benign
- Speech and object recognition
- Robot navigation

# Machine Learning and Big Data

- Many machine learning algorithms are data-driven
- Data provides the “experience” that ML algorithms use to train
- More data generally leads to better outcomes
  - Not always true
- ML and big data go hand in hand

# Three Primary Categories

- Supervised Learning

- Provide sample inputs and the desired outputs (a.k.a. **labels**)
- Algorithm learns function  $f$  that maps inputs to outputs

- Unsupervised Learning

- Provide only data (no labels)
- Algorithm figures out underlying structure

- Reinforcement Learning

- Computer is placed in a dynamic environment
- Learns how to act in order to maximize **reward** or minimize **penalty**

# Machine Learning vs Traditional Programming

- Traditional programming:



- Supervised machine learning:



- Unsupervised machine learning:



# Supervised Learning



# Supervised vs Unsupervised Learning

- Unsupervised learning → no notion of “right answer”
  - Many possible answers
- Supervised learning → there is a right answer!
  - Algorithms are trained (“supervised”) with training data
  - Data often divided into training and test data sets
    - Supervised learning on training set
    - Model verification on test set

# Supervised Learning Goal

- The ultimate goal is to develop a “model”,  $P(x)$ 
  - “ $x$ ” describes the **features** of the input
- We use training data to develop  $P(x)$
- Then, we can use  $P(x)$  to make “predictions” on new data

# Supervised Learning Goal

- Example: train a predictor  $P(x)$  that predicts house price based on:
  - Size (sq. ft.)
  - GPS coordinates
  - Number of bedrooms
  - Year built
- In this case, “ $x$ ” = {size, # bedrooms, year}
- In reality, the number of features could be millions!
  - Feature selection is a major challenge in many cases

# Types of Supervised Learning

- There are two major categories of supervised learning:
- Regression
  - For systems where the values are “continuous”
  - “Given a size and GPS coordinates, model how much this house costs”
- Classification
  - For systems where the output labels are discrete
  - “Is this cell benign or cancerous?” “male or female”


# Regression

# Regression Learning

- Suited for systems where at least some of the parameters are continuous
- Examples: modeling house prices, financial modeling
- In this case,  $P(x)$  may be a (piecewise) continuous function

# Fitting a Line

- Simplest form of regression learning
- Find a line that best fits the data




# Linear Regression Learning Goal

- Predictor  $P(x)$  has the form of a line:
- $P(x) = \theta_0 + \theta_1 x$
- Goal: Find  $\theta_0$  and  $\theta_1$  that best match the data



# Polynomial Relationship

- Sometimes the data is not strictly linear



# EE351K: Estimation

Remember point estimation we did in EE351K

1. Finding least squares estimator (LSE)?
2. Regression & Estimation related
  - a. Find error function (squared, example)
  - b. Find curve that minimizes error function
  - c. Depending on error function, different curve (potentially)

# Classification

# Classification Learning

- Suited for discrete labeling.
- “Is this cell cancerous?”: Binary labeling
- In this case,  $P(\mathbf{x})$  can be a likelihood of true/false (probabilistic classifier)
  - $P(\mathbf{x}) = 0 \Rightarrow$  Cell is definitely not cancerous
  - $P(\mathbf{x}) = 1 \Rightarrow$  Cell is definitely cancerous
  - $P(\mathbf{x}) = 0.7 \Rightarrow$  Cell is probably cancerous
  - $P(\mathbf{x}) = 0.5 \Rightarrow$  I don't know

# EE351K: Hypothesis Testing

Remember hypothesis testing from 351K?

Multiple hypotheses + Training data


Probabilistic classification = find most likely hypothesis

Classification has many other techniques (non probabilistic)

# Classification Learning: Cookie Example

Red - Bad cookie

Blue - Good cookie




Example from:  
<http://www.toptal.com/machine-learning/machine-learning-theory-an-introductory->

# Classification Learning: Cookie Example

Red - Bad cookie

Blue - Good cookie



Decision surface





Example from:  
<http://www.toptal.com/machine-learning/machine-learning-theory-an-introductory->

# Classification Learning: Linear Separability

- Linearly separable:



- Not linearly separable:





# Classification Learning: The Perceptron

- A perceptron is a **binary** classifier
- Classifies input  $\mathbf{x}$  into a single binary value

$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

- $w$  is a vector of weights
- $b$  is the y-intercept or **bias**

# Classification Learning: The Perceptron

- Training the perceptron involves determining the weight vector and the bias.
- The perceptron can only learn linearly separable patterns
- This is a big shortcoming, which we will address later

# Unsupervised Learning


# Unsupervised Learning Approaches

- Remember:
  - No “right answer”
  - Algorithm must discover underlying structure
- Approaches:
  - Clustering
  - Method of moments
  - Blind signal separation
    - PCA/ICA
    - Other matrix decomposition type techniques

# Clustering


Data points broken into  
“clusters”

Separated here by affine  
“dividers”



# Clustering


- Group a set of objects into clusters
- “Similar” objects should be in the same cluster



How many  
clusterings?


# Clustering by Position

- Total of three clusters



# Clustering by Color

- Three clusters again, but different points!
- Note the complexity of the fit





# Clustering Applications

- Marketing
  - Find distinct groups in your customer base (men, women, children, etc)
  - Develop products suited for each group
- Cell tower deployment
  - Cluster customers by location
  - Place towers at cluster centers to maximize QoS for customers
- Many more

# Method of Moments

- $E[x^k]$  = kth moment of  $x$
- Moments tell you a lot about the underlying distribution
- If  $X$  is a vector, its kth moment is a kth-order tensor
- Each value tells you interdependencies (or lack thereof) between components of  $X$ .


# Blind Signal Separation

- $N$  unknown sources  $s_j$
- One unknown operator  $\mathbf{A}$
- $P$  observed signals  $x_i$  with the relation:  $x = \mathbf{A}(s)$
- Goal: Estimate  $s$

# Blind Signal Separation: Cocktail Party Problem

- There are multiple people talking at a cocktail party
- You have recordings of the overall sound signature from multiple microphones
  - At least as many microphones as people
- Can you recover the individual voices of the people?

# Blind Signal Separation: Find $Z$



# Blind Signal Separation: Algorithms

- Need a way of separating **independent** sources
- Find a statistical representation of the data
- Project data onto a new set of axes that fulfill some statistical criteria
- Principal Component Analysis (PCA)
  - use an orthogonal transformation to generate linearly uncorrelated values
- Independent Component Analysis (ICA)
  - Not necessarily uncorrelated (for non-Gaussian sources)

# Reinforcement Learning

# Reinforcement Learning Approaches

- Reinforcement learning deals with how agents should act to maximize reward
  - Extremely general concept that applies to many fields
  - Also called Artificial Intelligence (AI)
- (Some) Approaches and Algorithms:
  - Artificial Neural networks
  - Simultaneous Localization and Mapping (SLAM)



# Neural Networks

# Neural Networks

- Family of modeling techniques inspired by biology
- Used to approximate functions with a large number of inputs
- Many different types:
  - Single perceptron
  - Feed-forward
  - Auto-encoders


# Neural Networks: Building Blocks

- Perceptrons are the building blocks of most neural nets
- A single layer of perceptrons can only learn linearly separable data
- But: Multiple layers of perceptrons are capable of learning more general patterns
  - Piecewise linear approximations

# Feed-Forward Networks

- Each node is a perceptron
- Output of previous layer is fed as input to the next layer

Example from:  
<http://www.toptal.com/machine-learning/an-introduction-to-deep-learning-from-perceptrons-to-deep-networks>



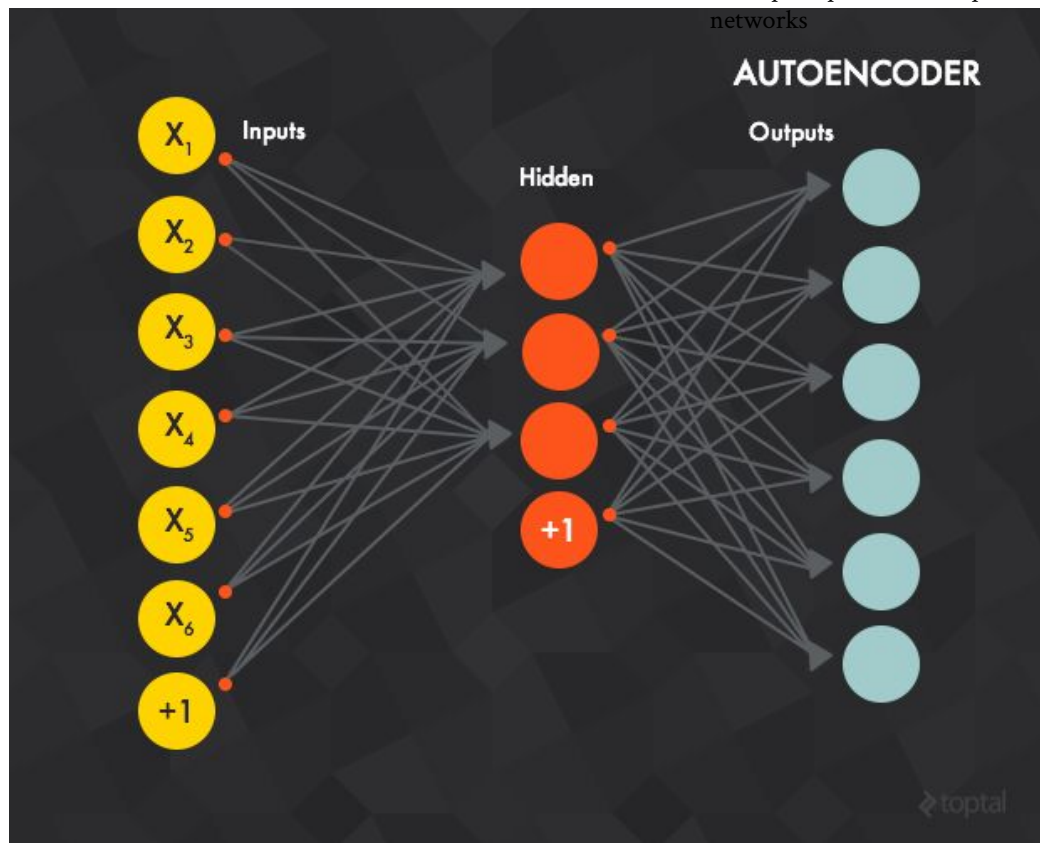
# Auto-Encoders

- A class of feed-forward neural nets
- Auto-encoders learn a compressed, distributed **encoding** (representation) of a dataset
- These networks recreate their input in a “compressed” form

# Auto-Encoders

- Fewer hidden nodes than inputs and outputs
- This forces the net to find a compressed representation of the input


Example from:  
<http://www.toptal.com/machine-learning/an-introduction-to-deep-learning-from-perceptrons-to-deep-networks>



# Deep Learning


- Auto-encoders can be viewed as “feature” detectors!
- What happens if we stack them?
- We get a net that progressively builds higher level features

Example from:  
<http://www.toptal.com/machine-learning/an-introduction-to-deep-learning-from-perceptrons-to-deep->



# Deep Learning: Hierarchical Feature Representation

Deep neural networks learn hierarchical feature representations





# Simultaneous Localization and Mapping (SLAM)

# SLAM

- Imagine a robot moving in an unknown environment
- Goal: we want to do two things (at the same time):
  - Figure out where we are in the environment
  - Generate a map of the environment
- But:
  - We don't know our **position** at any given time
  - We don't know our **velocity** at any given time
  - Our sensors are imperfect

# SLAM

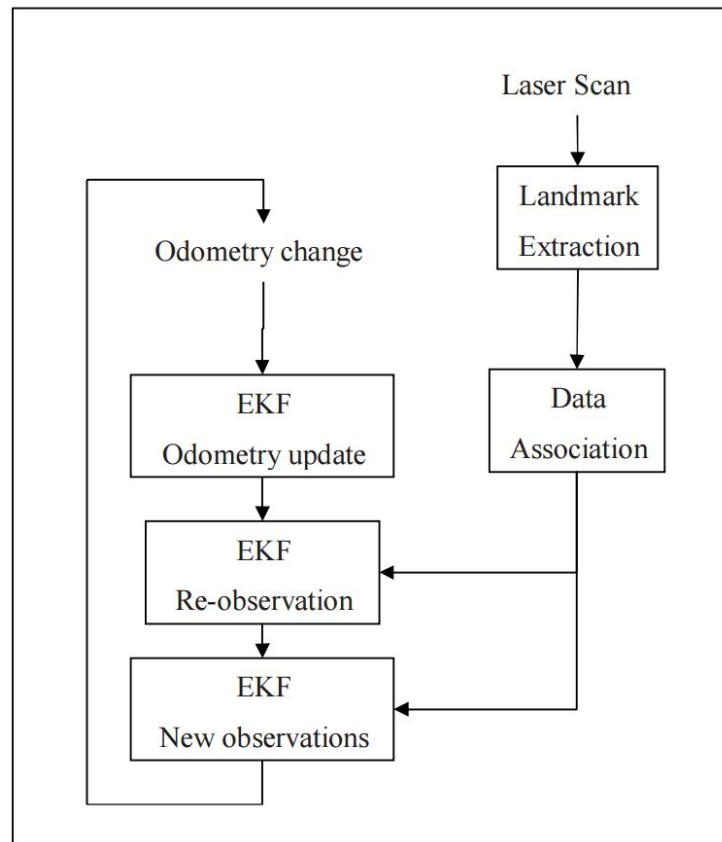
- Chicken and egg problem:
  - To build a map, we need to know our position
  - To know our position, we need a map
- Solution: Find both simultaneously
  - Localization: my position
  - Mapping: My environment

# SLAM: Required Components

- A robot
- Odometry for position and velocity data
  - An encoder on your motors
  - An airspeed sensor for flying craft
  - GPS
- Range measurement device (usually a laser/IR scanner)
- Other sensors: Acoustic, Visual (video) etc

# SLAM: Algorithm Flow

- **EKF** - Extended Kalman Filter
  - EKF helps the robot determine it's current position, attitude, and velocity
- **Landmark** - distinguishable features in the environment
  - These help the robot understand it's location and orientation in the environment



# SLAM

- Very complex problem
- Strong applications in robotics:
  - Self-driving vehicles
  - Path finding
  - Etc.
- Aerial SLAM

# Conclusion

- Machine learning is a very large field with many subcategories
  - Regression, classification, localization, etc
- There are three primary classes of machine learning algorithms
  - Supervised - learn from training data
  - Unsupervised - find underlying structures in data
  - Reinforcement - deal specifically with minimizing penalties
- Machine learning is a large part of big data analytics

# Introduction to K-means Clustering




# K-Means Clustering Algorithm

- Partitions  $N$  input points into  $K$  clusters
- Clusters defined by their **centroids** (“centers”)
- Each input point  $\mathbf{n}_i$  belongs to some cluster  $\mathbf{k}$  (with centroid  $\mathbf{c}_k$ )
- Minimize distance from each point to the centroid of its cluster


# K-Means Example

- Segment the following points into three clusters (by position)




# K-Means Example

1.  $K = 3$ . Pick three random data points as initial centroids




# K-Means Example

2. Assign each point to its nearest centroid



# K-Means Example

3. Recompute each centroid location as the average location of its children



# K-Means Example

4. Repeat assignment and update steps until no assignments change

